

DAIT Project

Trajectory Prediction for Human-Human Interaction

Rodolphe Farrando, Romain Gratier


23.05.2018

EPFL

Introduction


Introduction

- Trajectory prediction is crucial for improving autonomous vehicles behaviour



Introduction


- Trajectory prediction is crucial for improving autonomous vehicles behaviour
- Could avoid situations seen in the ethical lectures



Previous work Social LSTM : Human Trajectory Prediction in Crowded Spaces

In their project, they used different components to make the structure:


- One LSTM per pedestrian
- Social Pooling
- Prediction per frame



Previous work Social LSTM : Human Trajectory Prediction in Crowded Spaces


In their project, they used different components to make the structure:

- One LSTM per pedestrian
- Social Pooling
- Prediction per frame



In our project we only use:

- One CNN, or one LSTM
- Prediction per pedestrian




Data

Pre-processing

The preprocessing is divided in 4 steps:


1. We isolate each trajectory along with his interaction, that is the other trajectories that are around within the same frames



Pre-processing

The preprocessing is divided in 4 steps:


1. We isolate each trajectory along with his interaction, that is the other trajectories that are around within the same frames
2. We normalize the trajectories such that the first point is at $(0,0)$ and the second is at $(0, y_1)$



Pre-processing

The preprocessing is divided in 4 steps:


1. We isolate each trajectory along with his interaction, that is the other trajectories that are around within the same frames
2. We normalize the trajectories such that the first point is at $(0, 0)$ and the second is at $(0, y_1)$
3. We calculate axis velocities V_x and V_y



Pre-processing

The preprocessing is divided in 4 steps:


1. We isolate each trajectory along with his interaction, that is the other trajectories that are around within the same frames
2. We normalize the trajectories such that the first point is at $(0,0)$ and the second is at $(0, y_1)$
3. We calculate axis velocities V_x and V_y
4. For each frame, if there is a interacting pedestrian we add its coordinates and speed otherwise we add zeros



Pre-processing

The preprocessing is divided in 4 steps:

1. We isolate each trajectory along with his interaction, that is the other trajectories that are around within the same frames
2. We normalize the trajectories such that the first point is at $(0,0)$ and the second is at $(0, y_1)$
3. We calculate axis velocities V_x and V_y
4. For each frame, if there is a interacting pedestrian we add its coordinates and speed otherwise we add zeros
5. Data augmentation (flip and add noise) cf. last slide



Data

We have a file with:

- Pedestrians ID
- Frame number
- Twenty sets of x and y coordinates per pedestrian

Frame Number	ID	x	y	V_x	V_y
0	i	0	0	0	0
10	i	0	y	0	V_x
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots

Data

We have a file with:


- Pedestrians ID
- Frame number
- Twenty sets of x and y coordinates per pedestrian

Frame Number	ID	x	y	V_x	V_y
0	i	0	0	0	0
10	i	0	y	0	V_x
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots

We want :

- Train on the 10 first coordinates
- Predict the next 10

Example of trajectory to predict



Outputs structure

Finally our inputs have the following shape: $[10, N, 4 * N_{inter}]$, with

- 10: sequence length
- N : The number of data
- $4 * N_{inter}$: 4 (being the x and y coordinates and V_x and V_y velocities) times the number of pedestrians interacting with the one of interest.

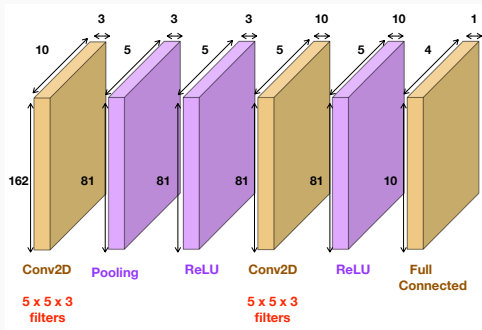
The four different cases are:

1. Predict coordinates with loss defines as $L_1 = (X - X_{pred})^2$ with $X = [x, y]$
2. Predict speeds with loss defines as $L_2 = (V - V_{pred})^2$ with $V = [V_x, V_y]$
3. Predict both coordinates and speeds with loss defines as $L = L_1 + L_2$
4. Predict both coordinates and speeds with loss defines as $L = L_1 + L_2 + L_3$, with $L_3 = (X - X_{t-1} + V_t * 0.4)^2$

The fourth case ensure that coordinates and speeds are not predicted independently.

Models


CNN



Inputs: sequence of coordinates and velocities of the pedestrians

Outputs: sequence of predicted coordinates and velocities for a trajectory of interest

LSTM



Inputs: sequence of coordinates and velocities of the trajectory of interest and of the interacting trajectories

Outputs: sequence of predicted coordinates and velocities for the trajectory of interest

Results

Results: Introduction

To calculate the correctness of the prediction two indicators are used:

1. The final displacement error: $e_{fin} = \sqrt{(X_n - X_{pred,n})^2}$
2. The mean displacement error: $e_{fin} = \sqrt{\frac{\sum_{i=0}^n (X_{gt,i} - X_{pred,i})^2}{(n)}}$

Depending on the inputs two ways are possible to find the predicted coordinates:

1. If the coordinates are predicted: directly use them
2. If the velocities are predicted: $X_t = X_{t-1} + V_t \cdot 0.4$, with 0.4 the time between two frames in seconds

Results for both models:

- 3 Types of trajectory defined: Static, linear and non-linear trajectories
- Mean and Final displacements

Results

Linear prediction results:

- Type 1: *Mean* = 0.141, *Final* = 0.322
- Type 2: *Mean* = 0.541, *Final* = 0.93
- Type 3: *Mean* = 0.651, *Final* = 1.457
- Total: *Mean* = 0.512, *Final* = 0.982

		CNN				LSTM			
		Type 1	Type 2	Type 3	Total	Type 1	Type 2	Type 3	Total
Coord.	Mean	4.696	5.144	4.674	4.176	1.309	0.777	0.862	0.877
	Final	10.246	7.009	10.501	5.602	1.385	0.92	1.108	1.037
Speed	Mean	0.567	5.133	1.911	4.17	0.726	0.573	0.651	0.616
	Final	0.77	6.971	3.882	5.587	1.412	1.045	1.231	1.148
2 Loss	Mean	1.269	5.134	1.762	4.163	0.695	0.532	0.627	0.581
	Final	2.727	6.978	3.546	5.57	1.302	0.963	1.2	1.076
3 Loss	Mean	0.549	5.135	3.829	4.163	0.748	0.607	0.681	0.647
	Final	0.758	6.983	4.962	5.573	1.364	1.072	1.308	1.177

LSTM and CNN comparison:

- For both LSTM and CNN, the "2 Loss" model have the better results
- For the CNN, results are much higher \Rightarrow the absence of memory is really penalising

LSTM and linear prediction comparison:


- Except for the first type trajectory, our "2 Loss" LSTM have better results
- Even for static trajectory, the LSTM tries to predict dynamic trajectory
- A few examples will show the difference

Representation

Representation

Examples of some trajectories:


Linear trajectories:



Representation

Examples of some trajectories:


Type 3 trajectory:



Representation

Examples of some trajectories:

We have also bad results:



Annex
